

$F2 \parallel WC_{\max}$ Structural Properties and Dynamic Programming

Author One^{a,*}, Author Two^b

^aInstitution One, Location, Country

^bInstitution Two, Location, Country

Abstract

The problem of minimizing maximum weighted completion time on a two-machine flow shop, denoted $F2 \parallel WC_{\max}$, is a scheduling problem we establish to be strongly NP-hard. We first establish that the optimal solution maintains the permutation property and introduce a novel block structure where jobs are partitioned by distinct weights, enabling efficient dynamic programming solutions. The exact dynamic programming algorithm leverages interval-based state representation and runs in $O(n \cdot K \cdot V^{3K-3})$ time for constant K , where V is the total processing time. For approximation guarantees, we design a fully polynomial-time approximation scheme (FPTAS) with state-class labeling that achieves $(1 + \varepsilon)$ -approximation while maintaining polynomial complexity. Additionally, we develop practical heuristics including an adapted NEH algorithm and an ant colony optimization (ACO) algorithm (Zhang et al., 2023) that provide fast solutions for large-scale instances. The computational results demonstrate that our algorithms efficiently solve both small and large instances while maintaining theoretical guarantees.

Keywords: Flow shop scheduling; Block-based dynamic programming; Approximation algorithms

1 Introduction

2 Preliminaries

We consider the problem $F2 \parallel WC_{\max}$. Let $\mathcal{J} = \{J_1, J_2, \dots, J_n\}$ be the set of n jobs to be processed on two machines M_1 and M_2 in series: each job J_j must be processed first on M_1 for

*Corresponding author. E-mail address: author@email.com

a_j time units and then on M_2 for b_j time units, with no preemption allowed. All processing times a_j and b_j are assumed to be positive integers. M_1 and M_2 can process at most one job at a time. Each job J_j has a weight $w_j > 0$. A schedule $\pi = (\pi(1), \pi(2), \dots, \pi(n))$ specifies the processing sequence on both machines. For a schedule π , let $A_j = \sum_{h=1}^j a_{\pi(h)}$ and $B_j = \sum_{h=1}^j b_{\pi(h)}$ be the cumulative M_1 and M_2 processing times after the first j jobs. The completion time of the job in position j on machine M_2 is

$$C_{\pi(j)} = \max_{1 \leq i \leq j} \left\{ \sum_{h=1}^i a_{\pi(h)} + \sum_{h=i}^j b_{\pi(h)} \right\}. \quad (1)$$

The objective is to find a schedule π that minimizes $WC_{\max}$.

The problem $F2 \parallel WC_{\max}$ can be formulated as a mixed integer linear program (MILP). For a permutation schedule, let the positions be indexed by $k = 1, \dots, n$. We introduce the following decision variables:

- $x_{jk} = 1$ if job J_j processing at position k , and 0 otherwise;
- C_k^1 completion time of the job at position k on M_1 ;
- C_k^2 completion time of the job at position k on M_2 ;
- C_j completion time of job J_j on M_2 ;
- Z the maximum weighted completion time $WC_{\max}$.
- $M = 2 \sum_{j=1}^n (a_j + b_j)$ be a sufficiently large constant.

The MILP model is:

$$\min \quad Z \quad (2)$$

$$\text{s.t.} \quad \sum_{k=1}^n x_{jk} = 1, \quad j = 1, \dots, n \quad (3)$$

$$\sum_{j=1}^n x_{jk} = 1, \quad k = 1, \dots, n \quad (4)$$

$$C_1^1 = \sum_{j=1}^n a_j x_{j1} \quad (5)$$

$$C_k^1 = C_{k-1}^1 + \sum_{j=1}^n a_j x_{jk}, \quad k = 2, \dots, n \quad (6)$$

$$C_1^2 = C_1^1 + \sum_{j=1}^n b_j x_{j1} \quad (7)$$

$$C_k^2 \geq C_k^1 + \sum_{j=1}^n b_j x_{jk}, \quad k = 2, \dots, n \quad (8)$$

$$C_k^2 \geq C_{k-1}^2 + \sum_{j=1}^n b_j x_{jk}, \quad k = 2, \dots, n \quad (9)$$

$$C_j \geq C_k^2 - M(1 - x_{jk}), \quad j = 1, \dots, n, \quad k = 1, \dots, n \quad (10)$$

$$Z \geq w_j C_j, \quad j = 1, \dots, n \quad (11)$$

$$x_{jk} \in \{0, 1\}, \quad (12)$$

$$C_k^1, C_k^2, C_j \geq 0 \quad (13)$$

Constraint (3)–(4) ensure a one-to-one assignment between jobs and positions. Constraints (5)–(6) compute the completion time of each position on M_1 . Constraints (7)–(9) compute the completion time of each position on M_2 : for $k \geq 2$, the job at position k starts on M_2 only after it completes M_1 and the preceding job releases M_2 . Constraint (10) links the position-based completion times to the job-based completion times via a big- M formulation: when $x_{jk} = 1$, job J_j is at position k and $C_j \geq C_k^2$; when $x_{jk} = 0$, the right-hand side becomes $C_k^2 - M \leq 0$, which is non-binding. Constraint (11) enforces $Z \geq w_j C_j$ for every job, so that minimizing Z yields the optimal $WC_{\max}$. Constraint (12) requires the assignment variables x_{jk} to be binary. Constraint (13) imposes the nonnegativity of the completion times C_k^1 , C_k^2 and C_j .

3 Problem formulation

Theorem 3.1. $F2 \parallel WC_{\max}$ is NP-hard.

Proof. We reduce from the Partition problem: given n positive integers $x_1, x_2, \dots, x_r$ with $\sum_{i=1}^r x_i = 2T$, does there exist subset Z_1 and $Z_2 \subseteq \{1, \dots, r\}$ such that: $\sum_{i \in Z_1} x_i = \sum_{i \in Z_2} x_i = T$?

Construct an instance of $F2 \parallel WC_{\max}$. Let the number of jobs $n = r + 1$. Let $a_j = 1$, $b_j = rx_j$, $w_j = T + 1$ for jobs $j = 1, \dots, n - 1$. Put $a_n = rT$, $b_n = 1$, $w_n = 2T + 1$ for job J_n . Does there exist a threshold Y such that:

$$WC_{\max} \leq Y = r(T + 1)(2T + 1)$$

holds?




Figure 3-1

We partition jobs j_1 to j_{n-1} into two subsets via j_n , the distribution of jobs on the machine is shown in Figure 3. Since the processing times of A_1 and A_2 are unknown, while their total processing time is $2T$. We may assume that the processing time of A_1 on M_2 is $r(T + \xi)$, A_2 is $r(T - \xi)$, where $-T \leq \xi \leq T$. Therefore, the completion time of J_n is determined by the following formula:

$$C_n^1 = |A_1| + rT + 1 \quad (14)$$

$$C_n^2 = 1 + r(T + \xi) + 1 \quad (15)$$

and $C_n = \max\{C_n^1, C_n^2\}$

If $\xi = 0$, then $\sum_{j \in A_1} b_j = \sum_{j \in A_2} b_j$, which means the partition problem has a feasible solution. In this case, the completion time of job J_n is determined by (14), we have:

$$\begin{aligned} C_n &= |A_1| + rT + 1 \leq (rT + r) = r(T + 1) \\ w_n C_n &= (2T + 1)(r(T + 1)) = Y \\ WC_{\max} &= r(T + 1)(C_n + rT) = Y \end{aligned}$$

so $w_j C_j \leq WC_{\max} = Y$

If the scheduling problem exists a threshold y such that $WC_{\max} \leq y$. When $\xi > 0$, C_n is determined by (15), we have:

$$\begin{aligned} C_n &= (1 + r(T + \xi) + 1) \geq (1 + r(T + 1) + 1) \\ w_n C_n &\geq (2T + 1)(1 + r(T + 1) + 1) = Y \end{aligned}$$

When $\xi < 0$, C_n is determined by (14), we have:

$$\begin{aligned} C_{\max} &= (C_n + \sum_{j \in A_2} b_j) \\ &= (1 + rT + |A_1| + r(T - \xi)) \\ &\geq (rT + r + rT) \\ &= r(2T + 1) = Y \end{aligned}$$

$$WC_{\max} \geq r(T+1)(r(2T+1)) = Y$$

so $WC_{\max} > Y$, there exists a contradiction.

Therefore, the Partition instance admits a solution if and only if the constructed $F2 \parallel WC_{\max}$ instance admits a schedule with $WC_{\max} \leq Y$. $\square$

Theorem 3.2. $F2 \parallel WC_{\max}$ is strongly NP-hard.

Proof. We reduce from the strongly NP-complete 3-Partition problem: given $3m$ positive integers $x_1, \dots, x_{3m}$ and an integer B satisfying $B/4 < x_i < B/2$ and $\sum_{i=1}^{3m} x_i = mB$, can the integers be partitioned into m triples, each of sum B ?

Let $R = 3m + 5$, $D = RB + 4$, and $T = mD$. We construct an instance with $3m$ element jobs $J_1, \dots, J_{3m}$ and m separator jobs $S_1, \dots, S_m$:

- for every element job J_j , set $a_j = 1$, $b_j = Rx_j + 1$;
- for every separator job S_k , set $a_{s_k} = RB + 1$, $b_{s_k} = 1$.

All processing times are positive integers. For $k = 1, \dots, m$, define $d_k = kD + 1$ and let $Q = (T + 2)^3$. Set

$$w_{s_k} = \left\lfloor \frac{Q}{d_k} \right\rfloor, \quad w_j = \left\lfloor \frac{Q}{T} \right\rfloor \quad (j = 1, \dots, 3m),$$

and use the decision threshold $Y = Q$. Since $d_k \leq T + 1$ and $Q \geq d_k(d_k + 1)$, we have $\lfloor Q/d_k \rfloor > d_k$. Thus, for every $d \in \{d_1, \dots, d_m, T\}$ and every integer completion time C ,

$$\left\lfloor \frac{Q}{d} \right\rfloor C \leq Q \implies C \leq d. \quad (16)$$

Indeed, if $C \geq d + 1$, write $Q = qd + r$ with $q = \lfloor Q/d \rfloor$ and $0 \leq r < d$. Then $q(d + 1) = Q - r + q > Q$ because $q > d > r$. Suppose first that the 3-Partition instance is a yes-instance. Let $G_1, \dots, G_m$ be triples with $\sum_{j \in G_k} x_j = B$ for every k , and schedule

$$G_1, S_1, G_2, S_2, \dots, G_m, S_m.$$

The three element jobs in G_k use 3 units on M_1 , and S_k uses $RB + 1$ units. Hence S_k completes on M_1 at time kD . The three element jobs in G_k use $RB + 3$ units on M_2 ; since the preceding separator completes on M_2 at time $(k - 1)D + 1$, the last element of G_k completes at time kD . Therefore

$$C_{s_k} = kD + 1 = d_k.$$

Every element job completes on M_2 no later than $T = mD$. By the definition of the weights, all weighted completion times are at most Q , so $WC_{\max} \leq Y$.

Conversely, suppose that a schedule π satisfies $WC_{\max} \leq Q$. By the standard permutation-schedule property for two-machine flow shops with a regular objective, it suffices to consider permutation schedules. The separator jobs have identical processing times. Therefore, if their positions are fixed, exchanging the labels of two separators that occur in the wrong order can only decrease the maximum weighted completion time, because $w_{s_1} \geq w_{s_2} \geq \dots \geq w_{s_m}$. We may consequently assume that the separators occur in the sequence $S_1, S_2, \dots, S_m$. By (16),

$$C_{s_k} \leq kD + 1 \quad (k = 1, \dots, m), \quad C_j \leq T \quad (j = 1, \dots, 3m).$$

No element job can occur after S_m . Indeed, if at least one element job did so, consider the last such element job. By the permutation-schedule property invoked above, the processing sequence on M_1 and M_2 coincides, so this job is also the last job on M_1 ; hence all m separators and all $3m$ element jobs have completed on M_1 no later than this job, and its completion on M_1 is at least the total M_1 processing time

$$m(RB + 1) + 3m = m(RB + 4) = T.$$

Its completion on M_2 would consequently be strictly greater than T , contradicting $C_j \leq T$. Thus the element jobs are partitioned into the m groups G_k lying between S_{k-1} and S_k , where S_0 denotes the beginning of the schedule. Put

$$B_k = \sum_{J_j \in G_k} x_j, \quad E_k = \sum_{h=1}^k |G_h|, \quad E_0 = 0.$$

For $k = 2, \dots, m$, the M_2 path gives

$$C_{s_k} \geq C_{s_{k-1}} + RB_k + |G_k| + 1,$$

while the M_1 path to S_{k-1} gives

$$C_{s_{k-1}} \geq (k-1)(RB + 1) + E_{k-1} + 1.$$

Combining these inequalities with $C_{s_k} \leq kD + 1$ yields

$$R(B_k - B) + E_k \leq 3k, \quad k = 2, \dots, m.$$

For $k = 1$, the M_2 path and $C_{s_1} \leq D + 1$ similarly give

$$R(B_1 - B) + E_1 \leq 4.$$

Because $R = 3m + 5$, these inequalities imply $B_k \leq B$ for every k : otherwise $B_k \geq B + 1$ and the respective left-hand side is greater than $3k$ (or greater than 4 when $k = 1$). All element jobs belong to one of the groups, so $\sum_{k=1}^m B_k = mB$. Hence $B_k = B$ for every k . Finally, $B/4 < x_j < B/2$ implies that a group with sum B contains exactly three elements: at most

two elements have sum strictly less than B , while at least four elements have sum strictly greater than B . Thus $G_1, \dots, G_m$ is a valid 3-partition.

The construction has $4m$ jobs, and all processing times, weights, and the threshold are bounded by a polynomial in m and B . Since 3-Partition remains NP-complete when B is polynomially bounded in m , this is a strong polynomial reduction. Therefore $F2 \parallel WC_{\max}$ is strongly NP-hard. $\square$




Figure 3-2: 3-Partition reduction

4 Algorithm

In this section, we design a dynamic programming algorithm for $F2 \parallel WC_{\max}$ and analyse its running time.

4.1 The number of weights is constant

For $F2 \parallel WC_{\max}$, an optimal schedule exists in which both machines process the jobs in the same sequence (Baker, 1974); we therefore restrict attention to such permutation schedules. Johnson's rule (Johnson, 1954) orders the n jobs so as to minimize the makespan: jobs with $a_j \leq b_j$ are scheduled first in non-decreasing sequence of a_j , and the remaining jobs follow in non-increasing sequence of b_j .

For a given schedule π , Let $W_1 > W_2 > \dots > W_K$ be the K distinct weights occurring among the jobs, and define $\mathcal{G}_k = \{J_j \in \mathcal{J} \mid w_j = W_k\}$, $k = 1, \dots, K$. For $k = 1, \dots, K$ in sequence, let ℓ_k be the position of the last job of weight W_k among the jobs not yet assigned to a block, i.e., among $\pi(\ell_{k-1} + 1), \dots, \pi(n)$, with the convention $\ell_k = \ell_{k-1}$ if no such job remains; set $\ell_0 = 0$. The blocks are then the segments

$$\mathcal{B}_k = \{\pi(\ell_{k-1} + 1), \dots, \pi(\ell_k)\}, \quad k = 1, \dots, K,$$

in particular, $\mathcal{B}_1 = \{\pi(1), \dots, \pi(\ell_1)\}$. Equivalently, after the blocks $\mathcal{B}_1, \dots, \mathcal{B}_k$ are removed, the remaining jobs are

$$\mathcal{J} \setminus \bigcup_{i=1}^k \mathcal{B}_i = \{\pi(\ell_k + 1), \dots, \pi(n)\}, \quad k = 1, \dots, K,$$

with $\ell_0 = 0$; consequently $\ell_1 \leq \ell_2 \leq \dots \leq \ell_K$ and no block extends past its own last job. For every block $\mathcal{B}_i$, let A_i denote the total processing time of its jobs on machine M_1 , B_i their total processing time on machine M_2 , C_i the completion time of its last job on machine M_2 , and L_i the weight of its last job. If $\ell_k = \ell_{k-1}$, then $\mathcal{B}_k$ contains no job and is called an empty block; an empty block $\mathcal{B}_i$ has $L_i = 0$ and $A_i = B_i = C_i = 0$, and may host any job of weight at most W_i . By construction, the last job of $\mathcal{B}_k$ carries weight W_k , so the block weights are non-increasing: $W_{\mathcal{B}_1} \geq W_{\mathcal{B}_2} \geq \dots \geq W_{\mathcal{B}_K}$.

We illustrate the block construction with a small example.

Example 4.1. Consider a 6-job instance with sequence $\pi = (J_1, \dots, J_6)$ and the following (a_j, b_j, w_j) values:

	J_1	J_2	J_3	J_4	J_5	J_6
a_j	2	5	6	8	4	3
b_j	5	7	9	4	3	2
w_j	2	2	5	2	4	2

The resulting groups are $\mathcal{G}_1 = \{J_3\}$, $\mathcal{G}_2 = \{J_5\}$ and $\mathcal{G}_3 = \{J_1, J_2, J_4, J_6\}$, and the blocks are $\mathcal{B}_1 = \{J_1, J_2, J_3\}$, $\mathcal{B}_2 = \{J_4, J_5\}$ and $\mathcal{B}_3 = \{J_6\}$.




Figure 4.1

Lemma 4.1. For any schedule π , rescheduling the jobs within each block according to Johnson's rule Johnson, 1954 and concatenating the resulting block schedules yields a schedule π' with $WC_{\max}(\pi') \leq WC_{\max}(\pi)$.

Proof. By construction, the final job of block $\mathcal{B}_k$ carries weight $W(\mathcal{B}_k)$, and every job in $\mathcal{B}_k$ has weight at most $W(\mathcal{B}_k)$: a weight occurring in $\mathcal{B}_k$ either equals $W(\mathcal{B}_k)$, or it has already appeared in an earlier block with larger weight. Let $C_k(\sigma)$ denote the completion time of the last job of $\mathcal{B}_k$ under schedule σ .

We first show that, within each block, the maximum weighted completion time is attained by the last job. For any $J_j \in \mathcal{B}_k$, the last job of the block finishes no earlier than J_j , so $C_j(\sigma) \leq C_k(\sigma)$, and $w_j \leq W(\mathcal{B}_k)$ by the block construction. Hence $w_j C_j(\sigma) \leq W(\mathcal{B}_k) \cdot C_k(\sigma)$. Taking the maximum over $j \in \mathcal{B}_k$ yields $\max_{J_j \in \mathcal{B}_k} w_j C_j(\sigma) = W(\mathcal{B}_k) \cdot C_k(\sigma)$, and therefore

$$WC_{\max}(\sigma) = \max_{k=1, \dots, K} \{W(\mathcal{B}_k) \cdot C_k(\sigma)\}.$$

Thus the overall objective depends only on the last job of each block.

Now let π' be the schedule obtained by rescheduling the jobs within each $\mathcal{B}_k$ according to Johnson's rule, while preserving the order of the blocks. We prove $C_k(\pi') \leq C_k(\pi)$ for all k by induction on k .

Base case $k = 1$. Block $\mathcal{B}_1$ starts at time zero on both machines regardless of its internal ordering, so Johnson's rule gives $C_1(\pi') \leq C_1(\pi)$.

Inductive step. Assume $C_{k-1}(\pi') \leq C_{k-1}(\pi)$. On M_1 , the start time of $\mathcal{B}_k$ equals the total M_1 work of $\mathcal{B}_1, \dots, \mathcal{B}_{k-1}$, which is invariant under internal rescheduling, so $\mathcal{B}_k$ begins at the same instant on M_1 under both schedules. On M_2 , $\mathcal{B}_k$ starts at $\max\{\sum_{i=1}^k A_i, C_{k-1}\}$, where the first term is identical under π and π' , and the second is no larger under π' by the induction hypothesis. Hence $\mathcal{B}_k$ starts no later on M_2 under π' than under π . Together with the same M_1 start and the internally optimal Johnson scheduling, this gives $C_k(\pi') \leq C_k(\pi)$, completing the induction.

Finally,

$$WC_{\max}(\pi') = \max_k W(\mathcal{B}_k) \cdot C_k(\pi') \leq \max_k W(\mathcal{B}_k) \cdot C_k(\pi) = WC_{\max}(\pi),$$

which establishes the lemma. $\square$

4.2 Algorithm Properties and Dynamic Programming

We now develop a dynamic programming algorithm that processes jobs sequentially in Johnson sequence and assigns each job to one of the K blocks $\mathcal{B}_1, \dots, \mathcal{B}_K$.

Lemma 4.1 provides the structural basis for the dynamic program: given any assignment of jobs to the K blocks, sequencing the jobs within each block $\mathcal{B}_i$ according to Johnson's rule and concatenating the blocks in non-increasing weight sequence yields a schedule whose $WC_{\max}$ value is no larger than that of any other schedule respecting the same assignment. For each job J_j ($j = 1, \dots, n$), let $\mathcal{H}_j$ denote the set of states attainable after the first j jobs have been processed. A state records, for every block $\mathcal{B}_i$:

- A_i : the total processing time on machine M_1 of the jobs in block $\mathcal{B}_i$;
- B_i : the total processing time on machine M_2 of the jobs in block $\mathcal{B}_i$;
- C_i : the completion time of the last job of block $\mathcal{B}_i$ on machine M_2 ;
- L_i : the weight of the last job of block $\mathcal{B}_i$.

Since $W_1 > \dots > W_K$, the blocks eligible for J_j form a prefix:

$$\mathcal{E}(J_j) = \{i : w_j \leq W_i\}, \quad i^* = \max_{i \in \{1, \dots, K\}} \{i : W_i \geq w_j\},$$

so that $\mathcal{E}(J_j) = \{1, \dots, i^*\}$. A state is in fact the aggregate K -block tuple, described by $4K - 2$ parameters: A_K and B_K need not be stored, because after the first j jobs have been processed, each of these jobs belongs to exactly one block, so the totals $\sum_{i=1}^K A_i = \sum_{h=1}^j a_h$ and $\sum_{i=1}^K B_i = \sum_{h=1}^j b_h$ are known, and A_K and B_K are recovered by

$$A_K = \sum_{h=1}^j a_h - \sum_{i=1}^{K-1} A_i, \quad B_K = \sum_{h=1}^j b_h - \sum_{i=1}^{K-1} B_i.$$

Hence each state of $\mathcal{H}_j$ is represented by the tuple

$$(A_1, \dots, A_{K-1}; B_1, \dots, B_{K-1}; C_1, \dots, C_K; L_1, \dots, L_K).$$

The initial state set contains a single state with all components set to zero: $\mathcal{H}_0 = \{(0, \dots, 0; 0, \dots, 0; 0, \dots, 0; 0, \dots, 0)\}$. In Algorithm 1 below, T_{A_i} denotes the total M_1 processing time of the first i blocks, and T_{B_i} the completion time of block $\mathcal{B}_i$ on M_2 .

Lemma 4.2. *Let $\mathcal{H} = (A_1, \dots, A_{K-1}; B_1, \dots, B_{K-1}; C_1, \dots, C_K; L_1, \dots, L_K)$ and $\mathcal{H}' = (A'_1, \dots, A'_{K-1}; B'_1, \dots, B'_{K-1}; C'_1, \dots, C'_K; L'_1, \dots, L'_K)$ be two states in $\mathcal{H}_j$ for job J_j . If $A_i = A'_i$, $B_i = B'_i$, $L_i = L'_i$ for all i , and $C_i \leq C'_i$ for all i , then $\mathcal{H}'$ can be discarded without affecting the optimal solution.*

Proof. Consider any continuation that extends state $\mathcal{H}'$ to a complete schedule Π' . Replacing the prefix represented by $\mathcal{H}'$ with the prefix represented by $\mathcal{H}$ yields a schedule Π that uses the same assignment of the remaining jobs. For every subsequent block $\mathcal{B}_k$ ($k \geq 2$), the M_2 start time is $\max\{\sum_{i=1}^k A_i, C_{k-1}\}$, where C_{k-1} is no larger under $\mathcal{H}$ than under $\mathcal{H}'$. By induction, the M_2 completion time of every subsequent block is no larger under $\mathcal{H}$, and since A_i , B_i and L_i are identical, the objective value $WC_{\max}$ of Π does not exceed that of Π' . $\square$

For $j = 1, \dots, n$, the algorithm extends $\mathcal{H}_{j-1}$ to $\mathcal{H}_j$ as follows: for each state in $\mathcal{H}_{j-1}$ and each block $i \in \mathcal{E}(J_j)$, the algorithm places J_j at the end of $\mathcal{B}_i$. Only the i -th slice of the state is updated, while all other blocks remain unchanged:

$$C'_i = \max\{C_i + b_j, A_i + a_j + b_j\}, \quad A'_i = A_i + a_j, \quad B'_i = B_i + b_j. \quad (17)$$

The resulting candidate states are collected in temporary sets $\mathcal{Y}_1, \dots, \mathcal{Y}_K$. By Lemma 4.2, merging the candidates removes dominated and duplicate states and yields $\mathcal{H}_j$.

Algorithm 1: Block concatenation and evaluation

Input: State $\mathcal{H} = \{A_i, B_i, C_i, L_i\}$.

1 **Step 1.** [Initialization] Set $T_{A_0} \leftarrow 0, T_{B_0} \leftarrow 0, WC_{\max} \leftarrow 0$.

2 **Step 2.** [Evaluation] **for** $i = 1$ **to** K **do**

3 $T_{A_i} \leftarrow T_{A_{i-1}} + A_i$

4 $T_{B_i} \leftarrow \max\{T_{A_{i-1}} + C_i, T_{B_{i-1}} + B_i\}$

5 **if** $L_i > 0$ **then**

6 $WC_{\max} \leftarrow \max\{WC_{\max}, L_i \cdot T_{B_i}\}$

7 **end**

8 **end**

Output: $WC_{\max}$.

Algorithm 2: DP for $F2 \parallel WC_{\max}$

Input: Jobs $\mathcal{J}$ with (a_j, b_j, w_j) ; max block count K ; weights $W_1 > \dots > W_K$.

1 **Step 1.** [Preprocessing] Set $L_i \leftarrow 0$ for $i = 1, \dots, K$.

2 **Step 2.** [Initialization] Set $\mathcal{H}_i \leftarrow \{(0, 0, 0, 0)_{i=1}^K\}$.

3 **Step 3.** [Generation] **for** $j = 1$ **to** n **do**

4 $\mathcal{Y}_i \leftarrow \emptyset$ for $i = 1, \dots, K$

5 **for each** $(\{A_i, B_i, C_i, L_i\}_{i=1}^K) \in \mathcal{H}_{j-1}$ **do**

6 **for** $i = 1$ **to** K **do**

7 $A'_i \leftarrow A_i + a_j, B'_i \leftarrow B_i + b_j, L'_i \leftarrow w_j$

8 $C'_i \leftarrow \max\{C_i + b_j, A_i + a_j + b_j\}$

9 Add the updated state to $\mathcal{Y}_i$

10 **end**

11 [Merging] $\mathcal{H}_j \leftarrow \text{MERGE}(\mathcal{Y}_1, \dots, \mathcal{Y}_K)$

12 **end**

13 **end**

14 **Step 4.** [Feasibility check] Remove from $\mathcal{H}_n = \{A_i, B_i, C_i, L_i\}_{i=1}^K$ violating the positional constraint $L_i \in \{0, W_i\}$.

15 **Step 5.** [Calculate] Set $Z^* \leftarrow +\infty$

16 **for each state** $S = \{A_i, B_i, C_i, L_i\}_{i=1}^K \in \mathcal{H}_n$ **do**

17 $WC_{\max} \leftarrow \text{Algorithm 1}((A_i, B_i, C_i, L_i))$

18 $Z^* \leftarrow \min\{Z^*, WC_{\max}\}$

19 **end**

20 **Step 6.** [Result] By standard backtracking Z^* .

Output: get the schedule σ .

To this end, each state in $\mathcal{H}_j$ stores a pointer to the state in $\mathcal{H}_{j-1}$ from which it was derived, together with the block $i \in \mathcal{E}(J_j)$ into which J_j was placed. Starting from the state in $\mathcal{H}_n$ attaining Z^* , backtracking along these pointers recovers, for every job J_j , the block $\mathcal{B}_i$ it belongs to. Sequencing the jobs of each block according to Johnson's rule (Lemma 4.1) and concatenating the blocks then yields an optimal schedule π^* with $WC_{\max}(\pi^*) = Z^*$.

Theorem 4.1. *Algorithm 2 solves $F2 \parallel WC_{\max}$ optimally in $O(n \cdot K \cdot |\mathcal{H}|_{\max})$ time, where $|\mathcal{H}|_{\max} \leq 2^K \cdot V^{3K-3}$, $V = \sum_{j=1}^n a_j + \sum_{j=1}^n b_j$. For constant K , this is $O(n \cdot V^{3K-3})$.*

Proof. Each state in $\mathcal{H}_j$ is a $(4K - 2)$ -tuple

$$(A_1, \dots, A_{K-1}; B_1, \dots, B_{K-1}; C_1, \dots, C_K; L_1, \dots, L_K).$$

We bound the number of distinct values each component can attain.

The quantity A_i is the total M_1 processing time of jobs assigned to block $\mathcal{B}_i$, hence $A_i \in [0, V]$. Since the j jobs under consideration form a prefix of the fixed Johnson sequence, $\sum_{i=1}^{K-1} A_i$ equals the sum of the a_h values over that prefix. Consequently at most $K - 1$ of the A_i are independent, contributing at most V^{K-1} distinct A -vectors. Symmetrically, $\sum_{i=1}^{K-1} B_i$ is the prefix sum of the b_j 's, so the B_i components likewise contribute at most V^{K-1} distinct combinations.

The quantity C_i is the makespan of $\mathcal{B}_i$ when scheduled in isolation by Johnson's rule. Since $C_i \leq A_i + B_i \leq V$ (each block's total M_1 and M_2 work is bounded by V), the C -vector contributes at most V^{K-1} distinct values. The weight L_i equals W_i if block $\mathcal{B}_i$ is nonempty and 0 if it is empty, yielding at most 2^K distinct L -vectors.

Multiplying these bounds yields

$$|\mathcal{H}_j| \leq 2^K \cdot V^{K-1} \cdot V^{2K-2} = 2^K \cdot V^{3K-3}.$$

In iteration j , the algorithm takes each state in $\mathcal{H}_{j-1}$, tries every block $i \in \{1, \dots, K\}$, computes the updated tuple in $O(1)$ time, and inserts it into $\mathcal{Y}_i$. The merge step collects the K sets $\mathcal{Y}_1, \dots, \mathcal{Y}_K$ and removes duplicates in time proportional to the total number of generated tuples. Hence iteration j costs $O(K \cdot |\mathcal{H}_{j-1}|)$, and summing over all n iterations gives $O(n \cdot K \cdot |\mathcal{H}|_{\max})$.

The final concatenation via Algorithm 1 evaluates each state in $\mathcal{H}_n$ in $O(K)$ time, which is dominated by the DP phase. When K is fixed, 2^K is constant, giving $O(n \cdot V^{3K-3})$. $\square$

4.3 A 2-Approximation Algorithm

The exact dynamic program of Section 4 runs in pseudo-polynomial time, and the FPTAS of Section 4.4 has running time exponential in the number K of distinct weights. When only a constant-factor guarantee is needed, the following elementary rule suffices.

Theorem 4.2. *Algorithm 3 runs in $O(n \log n)$ time and returns a schedule with $WC_{\max} \leq 2WC_{\max}^*$.*

Proof. Let σ be the schedule returned by Algorithm 3, let j_k be the job in position k of σ , and let $S_k = \{j_1, \dots, j_k\}$. Since the jobs are sorted by non-increasing weight, every job of S_k has weight at least w_{j_k} .

Algorithm 3: 2-Approximation for $F2 \parallel WC_{\max}$

Input: Jobs $\mathcal{J}$ with (a_j, b_j, w_j) .

1 Step 1. Sort the jobs by non-increasing weight w_j (ties broken arbitrarily).

2 Step 2. Schedule the jobs in this sequence on both machines.

Output: The resulting permutation schedule σ .

We first bound the completion times in σ . By (1), the job in position k completes on M_2 at time $C_{j_k}(\sigma)$, and each term in the maximum is bounded by the full sum:

$$\sum_{h=1}^i a_{j_h} + \sum_{h=i}^k b_{j_h} \leq \sum_{h=1}^k a_{j_h} + \sum_{h=1}^k b_{j_h} = \sum_{h \in S_k} (a_h + b_h), \quad 1 \leq i \leq k.$$

Hence

$$C_{j_k}(\sigma) \leq \sum_{h \in S_k} (a_h + b_h).$$

Since $WC_{\max}(\sigma) = \max_{1 \leq k \leq n} w_{j_k} C_{j_k}(\sigma)$, multiplying the inequality $C_{j_k}(\sigma) \leq \sum_{h \in S_k} (a_h + b_h)$ by $w_{j_k} > 0$ and taking the maximum over k gives

$$WC_{\max}(\sigma) = \max_{1 \leq k \leq n} w_{j_k} C_{j_k}(\sigma) \leq \max_{1 \leq k \leq n} w_{j_k} \sum_{h \in S_k} (a_h + b_h).$$

It remains to bound the right-hand side by $2WC_{\max}^*$. Fix a prefix S_k and let π^* be an optimal schedule. Let u be the job of S_k that finishes last on M_1 in π^* , and v the job of S_k that finishes last on M_2 . All jobs of S_k are processed on M_1 before u completes M_1 , so $C_u^* \geq \sum_{h \in S_k} a_h$; likewise, all M_2 processing of the jobs in S_k precedes v 's completion on M_2 , so $C_v^* \geq \sum_{h \in S_k} b_h$. Since $u, v \in S_k$, we have $w_u, w_v \geq w_{j_k}$. Because $WC_{\max}^*$ bounds the weighted completion time of every job in π^* , the bound $C_u^* \geq \sum_{h \in S_k} a_h$ together with $w_u \geq w_{j_k}$ gives

$$WC_{\max}^* \geq w_u C_u^* \geq w_{j_k} \sum_{h \in S_k} a_h,$$

and, similarly, $C_v^* \geq \sum_{h \in S_k} b_h$ together with $w_v \geq w_{j_k}$ gives

$$WC_{\max}^* \geq w_v C_v^* \geq w_{j_k} \sum_{h \in S_k} b_h.$$

Combining these two lower bounds yields

$$WC_{\max}^* \geq w_{j_k} \max \left\{ \sum_{h \in S_k} a_h, \sum_{h \in S_k} b_h \right\} \geq \frac{w_{j_k}}{2} \left(\sum_{h \in S_k} a_h + \sum_{h \in S_k} b_h \right),$$

where the last inequality uses $\max\{x, y\} \geq (x + y)/2$. Hence $w_{j_k} \sum_{h \in S_k} (a_h + b_h) \leq 2WC_{\max}^*$ for every k , and the claim follows from the upper bound above. $\square$

4.4 FPTAS

We now convert the exact dynamic program of Algorithm 2 into a fully polynomial-time approximation scheme by rounding the input data (Hall, 1994). One key point of this technique is the determination of valid lower and upper bounds for the coordinates to be rounded: each time coordinate ranges in a bounded interval, and rounding it to one of a fixed number of geometrically spaced values keeps the relative error within a single factor δ . The total processing times A_i, B_i of block $\mathcal{B}_i$ on M_1 and M_2 , and its completion time C_i , are rounded. The weight L_i of the last job of block $\mathcal{B}_i$ is not rounded; following Algorithm 2, the positional constraint $L_i \in \{0, W_i\}$ is required for every block, and a state violating it is discarded.

Fix $0 < \varepsilon < 1$ and let

$$\delta = 1 + \frac{\varepsilon}{2n}, \quad V_A = \sum_{j=1}^n a_j, \quad V_B = \sum_{j=1}^n b_j, \quad V = V_A + V_B.$$

For every block $\mathcal{B}_i$, the recurrence (17) preserves $A_i \leq V_A$ and $B_i \leq V_B$, since block $\mathcal{B}_i$ contains a subset of the jobs; moreover $C_i \leq A_i + B_i \leq V$. We therefore round the A dimension over $[0, V_A]$, the B dimension over $[0, V_B]$, and the C dimension over $[0, V]$. To this end, for $x \in \{A, B, C\}$ set $v^x = \lceil \log_\delta V_x \rceil$, where $V_C = V$, and split the range $[0, V_x]$ of dimension x into the $v^x + 1$ intervals

$$I_0^x = [0, 1), \quad I_r^x = [\delta^{r-1}, \delta^r) \quad (1 \leq r \leq v^x - 1), \quad I_{v^x}^x = [\delta^{v^x-1}, V_x].$$

All processing times are integral, so every value taken by a coordinate A_i, B_i, C_i is a non-negative integer. Hence any two values lying in a common interval satisfy $x \leq \delta y$; we call this the *interval-dominance* property, and it is the only fact about the rounding used in the analysis below.

For each block $\mathcal{B}_i$, let $r_i^C \in \{0, \dots, v^C\}$, $r_i^A \in \{0, \dots, v^A\}$ and $r_i^B \in \{0, \dots, v^B\}$ be the indices of the intervals containing C_i , A_i and B_i , respectively. The *label* of a state $\mathcal{H}^\Phi = (A_i, B_i, C_i, L_i)_{i=1}^K$ is the tuple

$$\ell(\mathcal{H}^\Phi) = (\mathbf{r}^A; \mathbf{r}^B; \mathbf{r}^C; \mathbf{L}),$$

where $\mathbf{r}^x = (r_1^x, \dots, r_K^x) \in \{0, \dots, v^x\}^K$ for $x \in \{A, B, C\}$, and $\mathbf{L} = (L_1, \dots, L_K)$. Two states share the same label exactly when, for every block, their coordinates A_i, B_i, C_i lie in the same geometric intervals and their block-end weights L_i coincide; that is, the two states are indistinguishable after rounding. There are $(v^A + 1)^K (v^B + 1)^K (v^C + 1)^K (K + 1)^K$ labels, a number independent of the job sizes; since $v^A, v^B \leq v^C$, this is at most $(v^C + 1)^{3K} (K + 1)^K$. The scheme keeps at most one state of each label: whenever several states carry the same label, all but one are deleted. In particular, states with the same label carry the same weight vector $\mathbf{L}$, so the retained state of an optimal trajectory has the same block-end weights as the trajectory itself. The weight L_i is the weight of the last job of block $\mathcal{B}_i$ and is not rounded.

Algorithm 4: State-class FPTAS for $F2 \parallel WC_{\max}$

Input: Jobs $\mathcal{J}$ with (a_j, b_j, w_j) ; block count K ; accuracy ε .

- 1 **Step 1.** [Preprocessing] Set $\delta = 1 + \varepsilon/(2n)$, $V_A = \sum_j a_j$, $V_B = \sum_j b_j$, $V = V_A + V_B$,
 $v^x = \lceil \log_\delta V_x \rceil$ for $x \in \{A, B, C\}$ with $V_C = V$, and I_r^x ($r = 0, \dots, v^x$).
- 2 **Step 2.** [Initialization] Set $\mathcal{H}^{\Phi,0} \leftarrow \{\mathcal{H}^{\Phi,0}\}$ with the zero state
 $\mathcal{H}^{\Phi,0} = (A_i, B_i, C_i, L_i)_{i=1}^K = (0, 0, 0, 0)_{i=1}^K$; set $L_i \leftarrow 0$ for $i = 1, \dots, K$.
- 3 **Step 3.** [Generation] **for** $j = 1$ **to** n **do**
- 4 $\mathcal{Y}_i \leftarrow \emptyset$ for $i = 1, \dots, K$
- 5 **for** each state $\mathcal{H}^\Phi = (A_i, B_i, C_i, L_i)_{i=1}^K \in \mathcal{H}_{j-1}^\Phi$ **do**
- 6 **for** $i = 1$ **to** K **do**
- 7 $\tilde{A}_i = A_i + a_j$, $\tilde{B}_i = B_i + b_j$, $\tilde{L}_i = w_j$
- 8 $\tilde{C}_i = \max\{C_i + b_j, A_i + a_j + b_j\}$
- 9 [Labeling] For $\tilde{A}_i, \tilde{B}_i, \tilde{C}_i$, determine the interval indices r_i^A, r_i^B, r_i^C such that
 $\tilde{A}_i \in I_{r_i^A}^A$, $\tilde{B}_i \in I_{r_i^B}^B$, $\tilde{C}_i \in I_{r_i^C}^C$. Record the label $\ell(\mathcal{H}^\Phi) = (\mathbf{r}^A; \mathbf{r}^B; \mathbf{r}^C; \mathbf{L})$ for
 the updated state.
- 10 **end**
- 11 [Deletion] From all candidate states in $\mathcal{Y}_{I_{r_1}}, \dots, \mathcal{Y}_{I_{r_K}}$ that share the same label,
 delete all but one; set $\mathcal{H}_j^\Phi$ to the remaining states.
- 12 **end**
- 13 **end**
- 14 **Step 4.** [Feasibility check] Same as Step 4 of Algorithm 2, applied to $\mathcal{H}^{\Phi,n}$.
- 15 **Step 5.** [Calculate] Same as Step 5 of Algorithm 2.
- 16 **Step 6.** [Result] Same as Step 6 of Algorithm 2.

Output: An approximate schedule σ .

Lemma 4.3. Let $\mathcal{H}^\Phi(j^*) = (A(j^*)_i, B(j^*)_i, C(j^*)_i, L(j^*)_i)_{i=1}^K$, $j = 0, \dots, n$, be the states along an optimal trajectory of Algorithm 2 with final value $WC_{\max}^*$. After processing j jobs, the state set $\mathcal{H}_j^\Phi$ of Algorithm 4 contains a state $\hat{\mathcal{H}}_j^\Phi = (\hat{A}_i, \hat{B}_i, \hat{C}_i, \hat{L}_i)_{i=1}^K$ with the same label $\ell(\hat{\mathcal{H}}_j^\Phi) = \ell(\mathcal{H}^\Phi(j^*))$ (in particular $\hat{L}_i = L(j^*)_i$ for every i) satisfying

$$\hat{A}_i \leq A(j^*)_i \delta^j, \quad \hat{B}_i \leq B(j^*)_i \delta^j, \quad \hat{C}_i \leq C(j^*)_i \delta^j \quad (1 \leq i \leq K). \quad (18)$$

Proof. We proceed by induction on j . The base case $j = 0$ is the zero state $\mathcal{H}^{\Phi,0}$ of Step 2 of Algorithm 4. Assume (18) holds for j , and let i^+ be the block to which $\mathcal{H}^\Phi(j^*)$ appends J_{j+1} in Algorithm 2. By the positional constraint of Algorithm 2, which assigns a job only to a block i with $w \leq W_i$, we have $w_{j+1} \leq W_{i^+}$; this constraint depends only on J_{j+1} and the block index, so the same transition is applicable from $\hat{\mathcal{H}}^{\Phi,j}$. Applying the exact recurrences to $\hat{\mathcal{H}}^{\Phi,j}$ and using $\delta^j \geq 1$ gives a candidate $\tilde{\mathcal{H}}^{\Phi,j+1}$ with

$$\begin{aligned} \tilde{A}_{i^+} &= \hat{A}_{i^+} + a_{j+1} \leq (A(j^*)_{i^+} + a_{j+1})\delta^j = A((j+1)^*)_{i^+}\delta^j, \\ \tilde{B}_{i^+} &= \hat{B}_{i^+} + b_{j+1} \leq (B(j^*)_{i^+} + b_{j+1})\delta^j = B((j+1)^*)_{i^+}\delta^j, \\ \tilde{C}_{i^+} &= \max\{\hat{C}_{i^+} + b_{j+1}, \hat{A}_{i^+} + a_{j+1} + b_{j+1}\} \end{aligned}$$

$$\leq \max \{C(j^*)_{i^+} + b_{j+1}, A(j^*)_{i^+} + a_{j+1} + b_{j+1}\} \delta^j = C((j+1)^*)_{i^+} \delta^j,$$

while the coordinates of the remaining blocks are unchanged and satisfy the bound at $j+1$. Moreover, since the optimal trajectory also appends J_{j+1} to i^+ , we have $\tilde{L}_{i^+} = w_{j+1} = L((j+1)^*)_{i^+}$, and the other blocks inherit $\hat{L}_i = L(j^*)_i$, so $\tilde{\mathcal{H}}^{\Phi, j+1}$ and $\mathcal{H}^{\Phi}((j+1)^*)$ carry the same label. The candidate $\tilde{\mathcal{H}}^{\Phi, j+1}$ carries the label $\ell(\tilde{\mathcal{H}}^{\Phi, j+1})$. By interval dominance, the state $\hat{\mathcal{H}}^{\Phi, j+1}$ kept in $\mathcal{H}^{\Phi, j+1}$ with that label satisfies $\hat{x}_i \leq \delta \tilde{x}_i$ for every $x \in \{A, B, C\}$ and every i ; this yields (18) at $j+1$ and completes the induction. Since every state of $\mathcal{H}^{\Phi, n}$ stores all $3K$ time coordinates, the induction covers the coordinates of the last block $\mathcal{B}_K$ as well. These coordinates are rounded and stored rather than recovered by prefix-sum subtraction as in the exact dynamic program, because the recovered values would accumulate the rounding errors of the first $K-1$ blocks and destroy the factor δ^j in (18). $\square$

Theorem 4.3. *For every $0 < \varepsilon < 1$, Algorithm 4 returns a schedule with*

$$WC_{\max} \leq (1 + \varepsilon) WC_{\max}^*,$$

where $WC_{\max}^*$ is the optimal value of F2 $\parallel WC_{\max}$.

Proof. Let $\mathcal{H}^{\Phi}(n^*)$ be the final state of the optimal trajectory of Algorithm 2 and $\hat{\mathcal{H}}^{\Phi, n}$ the state of $\mathcal{H}^{\Phi, n}$ guaranteed by Lemma 4.3 at $j = n$. By the lemma, $\hat{\mathcal{H}}^{\Phi, n}$ and $\mathcal{H}^{\Phi}(n^*)$ share the same label, hence $\hat{L}_i = L(n^*)_i$ for all i ; since $\mathcal{H}^{\Phi}(n^*)$ satisfies the positional constraint, $\hat{L}_i \in \{0, W_i\}$, and the nonempty blocks of $\hat{\mathcal{H}}^{\Phi, n}$ carry their weight W_i . In the block-concatenation evaluation of Algorithm 1, the quantity $T_{B, i}$ is nondecreasing in every coordinate and positively homogeneous, so $T_{B, i}(\hat{\mathcal{H}}^{\Phi, n}) \leq \delta^n T_{B, i}(\mathcal{H}^{\Phi}(n^*))$. Since the block weights W_i are fixed,

$$WC_{\max}(\hat{\mathcal{H}}^{\Phi, n}) = \max_{1 \leq i \leq K} W_i T_{B, i}(\hat{\mathcal{H}}^{\Phi, n}) \leq \delta^n \max_{1 \leq i \leq K} W_i T_{B, i}(\mathcal{H}^{\Phi}(n^*)) = \delta^n WC_{\max}^*.$$

Finally, $\delta^n = (1 + \varepsilon/(2n))^n \leq 1 + \varepsilon$ for $0 < \varepsilon < 1$. The schedule returned by Algorithm 4 is no worse than the schedule obtained by concatenating the blocks of the state $\hat{\mathcal{H}}^{\Phi, n}$, so its value is at most $(1 + \varepsilon) WC_{\max}^*$. $\square$

Theorem 4.4. *For fixed K , Algorithm 4 runs in*

$$O\left(n^{3K+1} \cdot K \cdot \frac{(\log V)^{3K}}{\varepsilon^{3K}}\right)$$

time and is therefore a fully polynomial-time approximation scheme for F2 $\parallel WC_{\max}$.

Proof. At each level j , Algorithm 4 keeps at most one state of each label, and there are $(v^A + 1)^K (v^B + 1)^K (v^C + 1)^K (K + 1)^K$ labels, at most $(v^C + 1)^{3K} (K + 1)^K$ since $v^A, v^B \leq v^C$. Each state generates K candidates in $O(1)$ time, and the final evaluation of the states of $\mathcal{H}^{\Phi, n}$ costs $O(K)$ per state. Hence the running time is

$$O(nK (v^C + 1)^{3K} (K + 1)^K),$$

which for constant K is $O(nK(v^C + 1)^{3K})$. Since $\ln(1 + \varepsilon/(2n)) \geq \varepsilon/(4n)$ for $0 < \varepsilon < 1$,

$$v^C = \left\lceil \frac{\ln V}{\ln \delta} \right\rceil = O\left(\frac{n \log V}{\varepsilon}\right),$$

and for constant K the bound reduces to $O(n^{3K+1}K(\log V)^{3K}/\varepsilon^{3K})$, polynomial in n , $\log V$, and $1/\varepsilon$. $\square$

Weight rounding. The FPTAS above rounds the time coordinates A_i, B_i, C_i and assumes that the number K of distinct weights is fixed. An analogous geometric rounding applies to the weights themselves when their ratio is bounded. Let $W_{\min} = \min_j w_j$, $W_{\max} = \max_j w_j$, and suppose $W_{\max}/W_{\min} \leq c$ for a constant c . Partition the interval $[W_{\min}, W_{\max}]$ into

$$J_0 = [W_{\min}, W_{\min}\delta), \quad J_r = [W_{\min}\delta^r, W_{\min}\delta^{r+1}) \quad (1 \leq r \leq q-1), \quad J_q = [W_{\min}\delta^q, W_{\max}],$$

with $q = \lceil \log_\delta(W_{\max}/W_{\min}) \rceil$, and round each weight w_j upward to the right endpoint of its interval,

$$\hat{w}_j = W_{\min} \delta^{r_j+1}, \quad r_j = \left\lfloor \log_\delta \frac{w_j}{W_{\min}} \right\rfloor,$$

so that $w_j \leq \hat{w}_j \leq \delta w_j$. The rounded weights take only $q+1 = O(\log_\delta(W_{\max}/W_{\min}))$ distinct values, so jobs whose weights differ by a factor at most δ merge into a single class. Since $\hat{w}_j \leq \delta w_j$ for every j , the optimal schedule $\hat{\sigma}$ of the rounded instance satisfies

$$WC_{\max}(\hat{\sigma}) \leq \hat{W}C_{\max}(\hat{\sigma}) = \hat{W}C_{\max}^* \leq \hat{W}C_{\max}(\pi^*) \leq \delta WC_{\max}^*,$$

where π^* is an optimal schedule of the original instance. Hence the weight rounding incurs a relative error of at most $\delta = 1 + \varepsilon/(2n) \leq 1 + \varepsilon$, exactly as in the completion-time rounding.

5 Heuristic Algorithms

The exact dynamic program of Section 4 is pseudo-polynomial, and the running time of the FPTAS of Section 4.4 is exponential in the number K of distinct weights. For large instances, fast heuristics are therefore of practical interest. In this section we present three heuristics for $F2 \parallel WC_{\max}$: an adaptation of the NEH algorithm (Puka et al., 2024), an ant colony optimization (ACO) algorithm (Zhang et al., 2023), and a genetic algorithm (Section 5.3).

Throughout this section, the objective of a permutation $\sigma = (\sigma(1), \dots, \sigma(k))$ of k jobs is evaluated by (1), so that

$$WC_{\max}(\sigma) = \max_{1 \leq t \leq k} w_{\sigma(t)} C_{\sigma(t)},$$

which is computed in $O(k)$ time.

5.1 NEH

The NEH algorithm (Puka et al., 2024) is one of the best-performing constructive heuristics for flow-shop scheduling. It orders the jobs by decreasing priority and then inserts them one by one into the partial schedule at the position that minimizes the objective.

For $F2 \parallel WC_{\max}$, jobs with large total processing time $a_j + b_j$ dominate the completion times and should be scheduled early. We therefore sort the jobs by non-increasing $a_j + b_j$, breaking ties by non-increasing weight w_j . The k -th job is inserted into each of the k possible positions of the current partial schedule, and a position attaining the smallest $WC_{\max}$ is kept.

Algorithm 5: NEH heuristic for $F2 \parallel WC_{\max}$

Input: Jobs $\mathcal{J}$ with (a_j, b_j, w_j) .

- 1 **Step 1.** [Sequence] Sort the jobs by non-increasing $a_j + b_j$, breaking ties by w_j ; let $J_1, \dots, J_n$ be the resulting sequence.
- 2 **Step 2.** [Initialization] $\sigma \leftarrow (J_1)$.
- 3 **Step 3.** [Insertion] **for** $k = 2$ **to** n **do**
- 4 $\sigma \leftarrow$ the sequence obtained by inserting J_k into σ at a position minimizing $WC_{\max}$.
- 5 **end**

Output: Schedule σ .

Theorem 5.1. *Algorithm 5 runs in $O(n^3)$ time.*

Proof. At step k , the job J_k is inserted into each of the k positions of a partial schedule of $k - 1$ jobs, and each candidate sequence is evaluated in $O(k)$ time by the recurrence above. Hence step k costs $O(k^2)$, and the total running time is $O(\sum_{k=1}^n k^2) = O(n^3)$. $\square$

5.2 Ant Colony Optimization

The ant colony optimization (ACO) algorithm (Zhang et al., 2023) is a metaheuristic inspired by the foraging behavior of ants. For $F2 \parallel WC_{\max}$, we adapt ACO to construct permutation schedules by iteratively improving solutions through pheromone trail reinforcement.

Algorithm 6: Ant Colony Optimization for $F2 \parallel WC_{\max}$

Input: Jobs $\mathcal{J}$ with (a_j, b_j, w_j) ; parameters α, β, ρ, Q ; number of ants m .

```
1 Step 1. [Initialization] Initialize pheromone matrix  $\tau_{ij} \leftarrow 1/n$  for all  $i, j$ ; set best solution  
    $\sigma^*, Z^* \leftarrow +\infty$ .  
2 Step 2. [Construction] for ant  $k = 1$  to  $m$  do  
3   Build a permutation  $\sigma_k$  by probabilistically selecting jobs:  
4   for position  $t = 1$  to  $n$  do  
5     Calculate probabilities  $p_{jt} = \frac{\tau_j^\alpha \cdot \eta_{jt}^\beta}{\sum_{l \in U} \tau_l^\alpha \cdot \eta_{lt}^\beta}$  where  $U$  is unselected jobs;  
6     Select job  $j$  with probability  $p_{jt}$  and set  $\sigma_k(t) = j$ ;  
7   end  
8   Evaluate  $Z_k = WC_{\max}(\sigma_k)$ ;  
9   if  $Z_k < Z^*$  then  
10     $\sigma^* \leftarrow \sigma_k, Z^* \leftarrow Z_k$   
11  end  
12 end  
13 Step 3. [Update Pheromone] for each edge  $(i, j)$  do  
14    $\tau_{ij} \leftarrow (1 - \rho) \cdot \tau_{ij} + \rho \cdot \Delta\tau_{ij}$  where  $\Delta\tau_{ij} = Q/Z^*$  if  $(i, j)$  in  $\sigma^*$ , else 0;  
15 end  
16 Step 4. [Termination] if stopping criterion not met then  
17   go to Step 2.  
18 end  
19  
Output: Best schedule  $\sigma^*$ .
```

5.3 A Genetic Algorithm

The objective $\max_j w_j C_j$ is a bottleneck criterion: the value of a schedule is determined by the single *critical* job that attains the maximum, and an ordering can be improved only by lowering that job's completion time, or by shortening completion times across the board. A genetic algorithm for $F2 \parallel WC_{\max}$ should therefore not treat all jobs symmetrically. We design its operators around the weight classes and the block structure of Section 4: high-weight jobs should occupy early positions (their weighted completion must be kept small), while the makespan should stay low so that no job is delayed unnecessarily.

Encoding and Decoding A chromosome is a permutation $\sigma = (\sigma(1), \dots, \sigma(n))$, where the gene at position t is the job scheduled t -th. Permutation encoding is adopted because it always represents a feasible permutation schedule: every job occurs exactly once, so no repair operator is needed to restore feasibility after crossover or mutation. To decode a chromosome, both machines process the jobs in the order $\sigma(1), \dots, \sigma(n)$; the completion time $C_{\sigma(t)}$ of the t -th job is obtained from the closed form (1), which is evaluated in $O(n)$

time using the prefix sums $A_t = \sum_{h=1}^t a_{\sigma(h)}$ and $B_t = \sum_{h=1}^t b_{\sigma(h)}$, and the objective value is

$$WC_{\max}(\sigma) = \max_{1 \leq t \leq n} w_{\sigma(t)} C_{\sigma(t)}.$$

The fitness $\text{fitness}(\sigma) = 1/WC_{\max}(\sigma)$ is used only to compare individuals, so its absolute scale is irrelevant.

Structured Initialization Random starting populations waste most of the search effort. The initial population P_0 of size N is therefore seeded with the three orderings that already realize the two structural ingredients of a good schedule—small makespan and early heavy jobs:

- the Johnson schedule σ^J produced by Johnson’s rule (Johnson, 1954), which minimizes the makespan and thereby tends to shorten every completion time;
- the weight-descending schedule σ^W , obtained by ordering jobs by non-increasing w_j with ties broken by Johnson’s rule, which directly attacks the bottleneck by giving heavy jobs the smallest possible completion times;
- the schedule σ^{NEH} returned by the NEH heuristic (Algorithm 5).

The remaining $N - 3$ individuals are uniform random permutations, which supply the diversity needed to escape these deterministic starting points and to explore regions of the search space that no constructive rule reaches. The population size is $N = \Theta(n)$: large enough to retain the block structure present in the seeds, yet small enough that a generation, whose cost is dominated by evaluating N individuals in $O(n)$ time each, remains $O(Nn)$.

Selection Binary tournament selection is used to form the mating pool. In each round, two individuals are drawn uniformly at random from the current population P_{g-1} , the fitter of the two wins and is copied into the mating pool M_g , and both are returned to the population (ties are broken arbitrarily); this is repeated until M_g contains N individuals. Tournament selection is preferred over roulette-wheel selection because it depends only on the *ranking* of fitness values rather than on their absolute magnitudes, so it remains selective even as the population converges and the fitness values of near-optimal schedules become very close.

Weight-Class Crossover (WCX) Crossover is the main source of new schedules, and it is here that the weight structure of Section 4 enters. Let $W_1 > \dots > W_K$ be the distinct weights and $\mathcal{G}_k = \{J_j \in \mathcal{J} \mid w_j = W_k\}$ the corresponding classes. A random threshold $q \in \{0, \dots, K\}$ fixes how many of the heaviest classes the child inherits from the first parent; the front of the child (the heavy jobs) is then copied from one parent while its tail (the light jobs) is filled from the other. This transmits the front-heavy layout shared by good

schedules of a min-max objective, while the order inside each part remains free to evolve. The threshold q is drawn uniformly so that every possible split point is explored with equal probability; the extreme cases $q = 0$ and $q = K$ reduce to copying σ_2 and σ_1 , respectively, so WCX degrades gracefully when the two parents coincide. Since the two parts are filled from different parents and each part preserves its parent’s relative order, the child is always a valid permutation and never requires repair.

Algorithm 7: Weight-Class Crossover (WCX) for $F2 \parallel WC_{\max}$

Input: Parents σ_1, σ_2 ; weight classes $\mathcal{G}_1, \dots, \mathcal{G}_K$.

- 1 **Step 1.** Draw q uniformly from $\{0, \dots, K\}$.
- 2 **Step 2.** Scan σ_1 and copy into σ , preserving order, every job that belongs to one of the q heaviest classes $\mathcal{G}_1, \dots, \mathcal{G}_q$.
- 3 **Step 3.** Scan σ_2 and append to σ , preserving order, every job not yet in σ (equivalently, every job of the classes $\mathcal{G}_{q+1}, \dots, \mathcal{G}_K$).

Output: Child schedule σ .

Critical-Job Mutation Mutation aims at the job that currently decides the objective. Let j^* be the critical job attaining $WC_{\max}(\sigma)$ at position t^* . Moving j^* to an earlier position $r < t^*$ can only decrease its own completion time (its machine-2 work is unchanged while it may start earlier), so the weighted completion $w_{j^*}C_{j^*}$ —currently the largest—tends to decrease. The mutation therefore tries every insertion of j^* into a position $r < t^*$, evaluates the resulting permutation by (1), and keeps the best; this is a cheap, targeted hill-climbing step. If no such insertion improves the value, a random swap of two jobs from different weight classes is performed with probability p_m to maintain diversity. Applied to every offspring, and combined with the elitist replacement, this operator acts as a local search that monotonically improves the incumbent without being trapped by the deterministic seeds.

Algorithm 8: Critical-Job Mutation for $F2 \parallel WC_{\max}$

Input: Schedule σ ; mutation probability p_m .

- 1 **Step 1.** [Identify] Evaluate $WC_{\max}(\sigma)$ by (1); let j^* be a job attaining the maximum and let t^* be its position in σ .
- 2 **Step 2.** [Insertion search] **for** $r = 1$ **to** $t^* - 1$ **do**
- 3 Form $\sigma^{(r)}$ by moving j^* to position r ; evaluate $WC_{\max}(\sigma^{(r)})$.
- 4 **end**
- 5 **Step 3.** **if** $\min_r WC_{\max}(\sigma^{(r)}) < WC_{\max}(\sigma)$ **then**
- 6 $\sigma \leftarrow \sigma^{(r^*)}$, where $r^* = \arg \min_r WC_{\max}(\sigma^{(r)})$
- 7 **else**
- 8 **if** $\text{rand}() < p_m$ **then**
- 9 Swap two jobs of σ lying in different weight classes.
- 10 **end**
- 11 **end**

Output: Improved schedule σ .

Genetic Algorithm Combining the ingredients above gives the following scheme. Elitism preserves the best λ solutions of each generation, and the final critical-job mutation plays the role of a local search.

Algorithm 9: Genetic Algorithm for $F2 \parallel WC_{\max}$

Input: Jobs $\mathcal{J}$; population size N ; generations G ; probabilities p_c, p_m ; elitism size λ .

- 1 **Step 1.** [Initialization] Build P_0 from the structured schedules $\sigma^J, \sigma^W, \sigma^{\text{NEH}}$ and $N - 3$ random permutations; evaluate all individuals; set σ^* to the best.
- 2 **Step 2.** [Evolution] **for** $g = 1$ **to** G **do**
- 3 **Selection** Form the mating pool M_g from P_{g-1} by binary tournament selection.
- 4 **Crossover** **while** $|O| < N$ **do**
- 5 Draw two parents from M_g ; with probability p_c apply Algorithm 7 to obtain two children, otherwise copy the parents into the offspring set O .
- 6 **end**
- 7 **Mutation** Apply Algorithm 8 to each offspring with probability p_m ; evaluate every child by (1).
- 8 **Replacement** $P_g \leftarrow$ the best λ individuals of P_{g-1} plus the $N - \lambda$ best offspring.
- 9 **Update** If the best individual of P_g improves on σ^* , replace σ^* by it.
- 10 **end**
- 11 **Step 3.** [Local search] Apply Algorithm 8 to σ^* until it yields no improvement.

Output: Best schedule σ^* .

Complexity Analysis Evaluating one schedule takes $O(n)$ time via (1); WCX costs $O(n)$; the critical-job mutation tries at most n insertions, each evaluated in $O(n)$ time, hence $O(n^2)$ per application. Initialization is dominated by the NEH seed, which costs $O(n^3)$ by

Theorem 5.1. A generation therefore costs $O(N n^2)$ in the worst case, giving $O(G N n^2)$ over G generations; for the typical setting $N, G = \Theta(n)$ the total running time is $O(n^4)$.

6 Conclusion

References

- Baker, K. R. (1974). *Introduction to Sequencing and Scheduling*. Wiley, New York.
- Hall, L. A. (1994). A polynomial approximation scheme for a constrained flow-shop scheduling problem. *Mathematics of Operations Research*, 19(1):68–85.
- Johnson, S. M. (1954). Optimal two- and three-stage production schedules with setup times included. *Naval Research Logistics Quarterly*, 1(1):61–68.
- Puka, R., Skalna, I., Duda, J., and Stawowy, A. (2024). Deterministic constructive vn-NEH+ algorithm to solve permutation flow shop scheduling problem with makespan criterion. *Computers & Operations Research*, 162:106473.
- Zhang, Y., Wang, X., and Li, Q. (2023). An improved ant colony optimization algorithm for flow shop scheduling problems with sequence-dependent setup times. *International Journal of Production Research*, 61(8):2456–2472.